

Heart Attack SHAP Practice (v2.1)

July 12, 2022

```
[1]: # This Heart Attack data Logistic Regression model is based on the 'heart.csv'
#      file (heart attack dataset from Kaggle)
# NOTE: Please do not take values as is, this is only a rough implementation of
#      → a LR ML model
#      Goal of this project is to show various ways SHAP can explain (and show)
#      → contribution
#      of each feature in the dataset
#      SHAP implementation is usually best for tree-based ML models such as
#      RandomForest, DecisionTree, etc

# More info on this dataset:

# Age : Age of the patient
# Sex : Sex of the patient
# cp : Chest Pain type
#      Value 0: typical angina
#      Value 1: atypical angina
#      Value 2: non-anginal pain
#      Value 3: asymptomatic

# trtbps : resting blood pressure (in mm Hg)
# chol: cholesterol in mg/dl fetched via BMI sensor
# fbs: (fasting blood sugar > 120 mg/dl)
#      1 = true
#      0 = false

# rest_ecg: resting electrocardiographic results
#      Value 0: normal
#      Value 1: having ST-T wave abnormality (T wave inversions and/or ST
#      → elevation or depression of > 0.05 mV)
#      Value 2: showing probable or definite left ventricular hypertrophy
#      → by Estes
```

```

# thalach: maximum heart rate achieved
# exang: exercise induced angina
#           1 = yes
#           0 = no

# old peak: ST depression induced by exercise relative to rest
# slp: the slope of the peak exercise ST segment
#           0 = unsloping
#           1 = flat
#           2 = downsloping

# caa: number of major vessels (0-3)
# thall : thalassemia
#           0 = null
#           1 = fixed defect
#           2 = normal
#           3 = reversable defect

# output: diagnosis of heart disease (angiographic disease status)
#           0: < 50% diameter narrowing. less chance of heart disease
#           1: > 50% diameter narrowing. more chance of heart disease

```

```

[2]: import os
import numpy as np
import pandas as pd
import matplotlib as plt
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix
import shap
import xgboost
import catboost as cb
from termcolor import colored as cl #text customization
from scipy.special import softmax
!export PATH=/Library/TeX/texbin:$PATH
shap.initjs() # load JS visualization code to notebook. SHAP plots won't be
↳displayed without this

```

<IPython.core.display.HTML object>

```

[3]: # Phase 1: import data

```

```
data = pd.read_csv('heart.csv')
data.head()

#print('\n', data.describe()) # this will show stats on each feature
```

```
[3]:
```

	age	sex	cp	trtbps	chol	fbs	restecg	thalachh	exng	oldpeak	slp	\
0	63	1	3	145	233	1	0	150	0	2.3	0	
1	37	1	2	130	250	0	1	187	0	3.5	0	
2	41	0	1	130	204	0	0	172	0	1.4	2	
3	56	1	1	120	236	0	1	178	0	0.8	2	
4	57	0	0	120	354	0	1	163	1	0.6	2	

	caa	thall	output
0	0	1	1
1	0	2	1
2	0	2	1
3	0	2	1
4	0	2	1

```
[4]: # Phase 2: Splitting dataset into X and y based on row and column using iloc[]
#      and reshaping X-variable before training and testing

X = data.iloc[:, :-1]
y = data.iloc[:, -1]

print('\nChecking X-variable before scaling\n', X)
print('\nChecking values in y\n', y)

# data partitioning
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25)

scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
print('\nChecking X-variable values after scaling\n', X_train)

# Checking training and testing samples for train_test_split
print('\nChecking data partitioning\nX_train samples: {}\nX_test samples: {}
→{}\nny_train samples: {}\nny_test samples: {}'.format(X_train, X_test,
→y_train, y_test))
```

Checking X-variable before scaling

	age	sex	cp	trtbps	chol	fbs	restecg	thalachh	exng	oldpeak	slp	\
--	-----	-----	----	--------	------	-----	---------	----------	------	---------	-----	---

0	63	1	3	145	233	1	0	150	0	2.3	0
1	37	1	2	130	250	0	1	187	0	3.5	0
2	41	0	1	130	204	0	0	172	0	1.4	2
3	56	1	1	120	236	0	1	178	0	0.8	2
4	57	0	0	120	354	0	1	163	1	0.6	2
..
298	57	0	0	140	241	0	1	123	1	0.2	1
299	45	1	3	110	264	0	1	132	0	1.2	1
300	68	1	0	144	193	1	1	141	0	3.4	1
301	57	1	0	130	131	0	1	115	1	1.2	1
302	57	0	1	130	236	0	0	174	0	0.0	1

	caa	thall
0	0	1
1	0	2
2	0	2
3	0	2
4	0	2
..
298	0	3
299	0	3
300	2	3
301	1	3
302	1	2

[303 rows x 13 columns]

Checking values in y

0	1
1	1
2	1
3	1
4	1
..	..
298	0
299	0
300	0
301	0
302	0

Name: output, Length: 303, dtype: int64

Checking X-variable values after scaling

```
[[ 1.2074225  0.68849544 -0.95607147 ...  0.99874214 -0.70333816
  1.10304227]
 [ 0.87070886 -1.45244244 -0.95607147 ... -0.64411631 -0.70333816
 -0.55517359]
 [ 0.30951945  0.68849544  0.9730683  ...  0.99874214 -0.70333816
 -0.55517359]
```

```

...
[ 0.87070886  0.68849544  0.00849841 ...  0.99874214 -0.70333816
-0.55517359]
[ 0.30951945 -1.45244244  0.00849841 ... -0.64411631  0.27018478
-0.55517359]
[ 0.98294674  0.68849544 -0.95607147 ...  0.99874214  2.21723067
1.10304227]]

```

Checking data partitioning

```

X_train samples: [[ 1.2074225  0.68849544 -0.95607147 ...  0.99874214
-0.70333816
1.10304227]
[ 0.87070886 -1.45244244 -0.95607147 ... -0.64411631 -0.70333816
-0.55517359]
[ 0.30951945  0.68849544  0.9730683 ...  0.99874214 -0.70333816
-0.55517359]
...
[ 0.87070886  0.68849544  0.00849841 ...  0.99874214 -0.70333816
-0.55517359]
[ 0.30951945 -1.45244244  0.00849841 ... -0.64411631  0.27018478
-0.55517359]
[ 0.98294674  0.68849544 -0.95607147 ...  0.99874214  2.21723067
1.10304227]]

```

```

X_test samples:      age  sex  cp  trtbps  chol  fbs  restecg  thalachh  exng
oldpeak  slp  \
209   59   1   0   140   177   0       1      162    1     0.0    2
5     57   1   0   140   192   0       1      148    0     0.4    1
299   45   1   3   110   264   0       1      132    0     1.2    1
274   47   1   0   110   275   0       0      118    1     1.0    1
59    57   0   0   128   303   0       0      159    0     0.0    2
..    ...   ...   ..   ...   ...   ...   ...   ...   ...   ...   ...
60    71   0   2   110   265   1       0      130    0     0.0    2
221   55   1   0   140   217   0       1      111    1     5.6    0
199   65   1   0   110   248   0       0      158    0     0.6    2
197   67   1   0   125   254   1       1      163    0     0.2    1
257   50   1   0   144   200   0       0      126    1     0.9    1

```

```

      caa  thall
209     1     3
5       0     1
299     0     3
274     1     2
59      1     2
..    ...   ...
60     1     2
221     0     3
199     2     1

```

```
197    2    3
257    0    3
```

```
[76 rows x 13 columns]
```

```
y_train samples: 31    1
```

```
258    0
9      1
28     1
90     1
```

```
..
113    1
```

```
284    0
```

```
137    1
```

```
302    0
```

```
217    0
```

```
Name: output, Length: 227, dtype: int64
```

```
y_test samples: 209    0
```

```
5      1
```

```
299    0
```

```
274    0
```

```
59     1
```

```
..
```

```
60     1
```

```
221    0
```

```
199    0
```

```
197    0
```

```
257    0
```

```
Name: output, Length: 76, dtype: int64
```

```
[5]: # Phase 3: Creating the model
# Understanding LR parameters:
#         solver = a string ('liblinear' by default) that decides what
#         ↪ solver to use for fitting the model
#         multi-class = a string ('ovr' by default) decides the approach
#         ↪ to use for handling multiple classes\
#         'ovr' says to make the binary fit for each class
#         'multinomial' says to apply the multinomial loss
#         ↪ fit
#         C = a positive floating-point number (1.0 by default) that
#         ↪ defines the relative strength of regularization
#         random_state = an integer, an instance of numpy.RandomState, or
#         ↪ None (default) that defines what pseudo number

model = LogisticRegression(solver='liblinear', C=0.5, multi_class='ovr',
#         ↪ random_state=0)
```

```

model.fit(X_train, y_train)

print('\nChecking LR output string\n', cl(model, attrs = ['bold']))
print('\nChecking LR coefficients: \n', model.coef_)
print('\nChecking LR intercept: \n', model.intercept_)

```

Checking LR output string

```

LogisticRegression(C=0.5, multi_class='ovr', random_state=0,
solver='liblinear')

```

Checking LR coefficients:

```

[[ 0.04938564 -0.48975268  1.02199574 -0.35408865  0.05784901  0.05261735
  0.41202293  0.24570615 -0.5247825  -0.60006623  0.35173065 -0.67646986
 -0.43941142]]

```

Checking LR intercept:

```

[0.25759399]

```

[6]: *# Phase 4:Evaluating the model: returns the matrix of probabilities*

```

X_test = scaler.transform(X_test)
y_pred = model.predict(X_test)
confusion_matrix(y_test, y_pred)

```

```

[6]: array([[29, 10],
          [ 4, 33]])

```

```

[7]: train_score = model.score(X_train, y_train)
test_score = model.score(X_test, y_test)
print('\nTrain score is: \n', train_score)
print('\nTest score is: \n', test_score)

```

Train score is:

```

0.8502202643171806

```

Test score is:

```

0.8157894736842105

```

[8]: `print(classification_report(y_test, y_pred))`

	precision	recall	f1-score	support
0	0.88	0.74	0.81	39
1	0.77	0.89	0.82	37

accuracy			0.82	76
macro avg	0.82	0.82	0.82	76
weighted avg	0.82	0.82	0.82	76

```
[9]: # Phase 5: Creating SHAP explainer object
#      Creating feature_names array to store features to be able to use in
#      →shap plots/graphs

feature_names = []
for feature in X.columns.tolist():
    feature_names.append(feature)
print('\nChecking feature names in columns from X-variable: \n', feature_names)

# Because we are using a LogisticRegression model which is a type of linear
#      →model, we want to use
#      the .LinearExplainer() SHAP method
# Use train data for masker -- > based on usual train/test paradigm, where you
#      →train your model
#      (and explainer) on train data, and try to predict (and explain) your
#      →test data.
#      (troubleshoot: https://stackoverflow.com/questions/66560839/
#      →what-do-maskers-really-do-
#      in-shap-package-and-fit-them-to-train-or-test)
#masker = shap.maskers.Independent(data = X_train)
#explainer = shap.LinearExplainer(model, masker = masker)
# Uncomment masker and explainer w/masker arg, Comment .Explainer to see actual
#      →representation (not
#      necessarily accurate). Doing this will cause labels on SHAP plots/
#      →graphs to display
#      'Feature #' instead of actual feature name

explainer = shap.Explainer(model.predict, X_test, feature_names=feature_names)
shap_values = explainer(X_train)
```

Checking feature names in columns from X-variable:

```
['age', 'sex', 'cp', 'trtbps', 'chol', 'fbs', 'restecg', 'thalachh', 'exng',
'oldpeak', 'slp', 'caa', 'thall']
```

```
[10]: # OPTIONAL: Take the mean average value of each feature to get feature
#      →importances determined by SHAP
```



```

# NOTE:      This is to only showcase how SHAP feature importance works (not my
↳own code, just put it
#
↳in here to provide comprehensive baseline of how SHAP
↳and feature importance
#
work)
# CREDIT:    Vinícius Trevisan (only this cell of the code)
#
↳https://towardsdatascience.com/
↳using-shap-values-to-explain-how-your-machine-learning-model-works-732b3f40e137

#def print_feature_importances_shap_values(shap_values, feature_names):
#    '''
#    Prints the feature importances based on SHAP values in an ordered way
#    shap_values -> The SHAP values calculated from a shap.Explainer object
#    features -> The name of the features, on the order presented to the
↳explainer
#    '''
# # Calculates the feature importance (mean absolute shap value) for each feature
# importances = []
# for i in range(shap_values.values.shape[1]):
#     importances.append(np.mean(np.abs(shap_values.values[:, i])))
# # Calculates the normalized version
# importances_norm = softmax(importances)
# # Organize the importances and columns in a dictionary
# feature_importances = {fea: imp for imp, fea in zip(importances,
↳feature_names)}
# feature_importances_norm = {fea: imp for imp, fea in zip(importances_norm,
↳feature_names)}
# # Sorts the dictionary
# feature_importances = {key:value for key, value in
↳sorted(feature_importances.items(), key=lambda item: item[1], reverse = True)}
# feature_importances_norm= {key:value for key,value in
↳sorted(feature_importances_norm.items(), key=lambda item: item[1], reverse =
↳True)}
# # Prints the feature importances
# for key, value in feature_importances.items():
#     print(f"{key} -> {value:.4f} (softmax = {feature_importances_norm[key]:
↳.4f})")

# print_feature_importances_shap_values(shap_values, feature_names)

```

```

[11]: # Phase 5a: Finding 'best' way to represent dataset using SHAP visualizations
# NOTE: Please do not take values as is, this is only a rough implementation of
↳a LR model
#
Goal of this project is to show various ways SHAP can explain (and
↳show) contribution
#
of each feature in the dataset

```

```

# for-loop goes through shap_values which stores .base_values, .values, .data
→while also looping
#           through feature_names to assign respective shap_value to its
→corresponding feature name
# Break the inner loop to execute string based on length of the feature_names
→array or else inner-loop
#           will repeat string format all # of times of length in feature_names
# The shap_values variable will have three attributes: .values, .base_values and
→.data.
# The .data attribute is simply a copy of the input data, .base_values is the
→expected value of the
#           target, or the average target value of all the train data, and .
→values are the SHAP
#           values for each example.
# If we are only interested in the SHAP values, we can use the explainer.
→shap_values() method

# If we are only interested in the SHAP values, we can use the explainer.
→shap_values() method:
print('\nArray of SHAP values with added feature_names: ')
for name in feature_names:
    print('\n')
    for value in shap_values:
        print("{} data and values {}".format(name, value))
    break

# By default, a SHAP beeswarm plot will take the mean absolute value of each
→feature
#           over all the instances (rows) of the dataset
# Each point on the chart is one SHAP value for a prediction and feature
# Red color means higher value of a feature. Blue means lower value of a feature.
# We can get the general sense of features' directionality impact based on the
→distribution of
#           the red and blue dots
print('\nPrinting Beeswarm SHAP Plot: \n')
shap.plots.beeswarm(shap_values)

# In this chart, the x-axis stands for the SHAP value and the y-axis has all the
→features

```

```

print('\nPrinting Summary SHAP Bar Plot: \n')
shap.summary_plot(shap_values, feature_names, plot_type='bar')

# Interesting SHAP scatter plot summary of all SHAP values for each feature
print('\nPrinting Scatter SHAP Plot for All Features: \n')
shap.plots.scatter(shap_values[:, :])

# for-loop for SHAP scatter plots to show correlation between different features
print('\nPrinting Scatter SHAP Plot Between Features: \n')
for feature in feature_names:
    shap.plots.scatter(shap_values[:, feature], color=shap_values)

# For analysis of local, instance-wise effects, we can use the following plots
→ on single
#           observations using "shap_values[0]".
print('\nPrinting Bar SHAP Plot for Local, Instance-Wise Effects: \n')
shap.plots.bar(shap_values[0])

# The waterfall plot has the same information, represented in a different manner.
→
# Here we can see how the sum of all the SHAP values equals the difference
→ between
#           the prediction  $f(x)$  and the expected value  $E[f(x)]$ .
print('\nPrinting Waterfall SHAP Plot for Local, Instance-Wise Effects: \n')
shap.plots.waterfall(shap_values[0])

# In this plot the positive SHAP values are displayed on the left side and the
→ negative on the
#           right side, as if competing against each other. The highlighted
→ value is the
#           prediction for that observation.
# Note that this is for a single instance
print('\nPrinting Force SHAP Plot for Local, Instance-Wise Effects: \n')
shap.plots.force(shap_values[0])

```

Array of SHAP values with added feature_names:

```

"age" data and values
.values =

```

```
array([ 0.01023392, -0.06688596, -0.13450292,  0.08479532, -0.00657895,  
       -0.00402047,  0.13596491, -0.01754386,  0.12426901,  0.08077485,  
       0.15899123,  0.16959064, -0.10087719])
```

```
.base_values =  
0.5657894736842105
```

```
.data =  
array([ 1.2074225 ,  0.68849544, -0.95607147, -0.64519349, -1.31346187,  
       -0.43414476,  0.93849623, -0.4464528 , -0.6815542 , -0.57895484,  
       0.99874214, -0.70333816,  1.10304227])
```

"sex" data and values

```
.values =  
array([ 0.01023392, -0.06688596, -0.13450292,  0.08479532, -0.00657895,  
       -0.00402047,  0.13596491, -0.01754386,  0.12426901,  0.08077485,  
       0.15899123,  0.16959064, -0.10087719])
```

```
.base_values =  
0.5657894736842105
```

```
.data =  
array([ 1.2074225 ,  0.68849544, -0.95607147, -0.64519349, -1.31346187,  
       -0.43414476,  0.93849623, -0.4464528 , -0.6815542 , -0.57895484,  
       0.99874214, -0.70333816,  1.10304227])
```

"cp" data and values

```
.values =  
array([ 0.01023392, -0.06688596, -0.13450292,  0.08479532, -0.00657895,  
       -0.00402047,  0.13596491, -0.01754386,  0.12426901,  0.08077485,  
       0.15899123,  0.16959064, -0.10087719])
```

```
.base_values =  
0.5657894736842105
```

```
.data =  
array([ 1.2074225 ,  0.68849544, -0.95607147, -0.64519349, -1.31346187,  
       -0.43414476,  0.93849623, -0.4464528 , -0.6815542 , -0.57895484,  
       0.99874214, -0.70333816,  1.10304227])
```

"trtbps" data and values

```
.values =  
array([ 0.01023392, -0.06688596, -0.13450292,  0.08479532, -0.00657895,  
       -0.00402047,  0.13596491, -0.01754386,  0.12426901,  0.08077485,  
       0.15899123,  0.16959064, -0.10087719])
```

```

.base_values =
0.5657894736842105

.data =
array([ 1.2074225 ,  0.68849544, -0.95607147, -0.64519349, -1.31346187,
        -0.43414476,  0.93849623, -0.4464528 , -0.6815542 , -0.57895484,
         0.99874214, -0.70333816,  1.10304227])

"chol" data and values
.values =
array([ 0.01023392, -0.06688596, -0.13450292,  0.08479532, -0.00657895,
        -0.00402047,  0.13596491, -0.01754386,  0.12426901,  0.08077485,
         0.15899123,  0.16959064, -0.10087719])

.base_values =
0.5657894736842105

.data =
array([ 1.2074225 ,  0.68849544, -0.95607147, -0.64519349, -1.31346187,
        -0.43414476,  0.93849623, -0.4464528 , -0.6815542 , -0.57895484,
         0.99874214, -0.70333816,  1.10304227])

"fbs" data and values
.values =
array([ 0.01023392, -0.06688596, -0.13450292,  0.08479532, -0.00657895,
        -0.00402047,  0.13596491, -0.01754386,  0.12426901,  0.08077485,
         0.15899123,  0.16959064, -0.10087719])

.base_values =
0.5657894736842105

.data =
array([ 1.2074225 ,  0.68849544, -0.95607147, -0.64519349, -1.31346187,
        -0.43414476,  0.93849623, -0.4464528 , -0.6815542 , -0.57895484,
         0.99874214, -0.70333816,  1.10304227])

"restecg" data and values
.values =
array([ 0.01023392, -0.06688596, -0.13450292,  0.08479532, -0.00657895,
        -0.00402047,  0.13596491, -0.01754386,  0.12426901,  0.08077485,
         0.15899123,  0.16959064, -0.10087719])

.base_values =
0.5657894736842105

```

```
.data =
array([ 1.2074225 ,  0.68849544, -0.95607147, -0.64519349, -1.31346187,
        -0.43414476,  0.93849623, -0.4464528 , -0.6815542 , -0.57895484,
         0.99874214, -0.70333816,  1.10304227])
```

"thalachh" data and values

```
.values =
array([ 0.01023392, -0.06688596, -0.13450292,  0.08479532, -0.00657895,
        -0.00402047,  0.13596491, -0.01754386,  0.12426901,  0.08077485,
         0.15899123,  0.16959064, -0.10087719])
```

```
.base_values =
0.5657894736842105
```

```
.data =
array([ 1.2074225 ,  0.68849544, -0.95607147, -0.64519349, -1.31346187,
        -0.43414476,  0.93849623, -0.4464528 , -0.6815542 , -0.57895484,
         0.99874214, -0.70333816,  1.10304227])
```

"exng" data and values

```
.values =
array([ 0.01023392, -0.06688596, -0.13450292,  0.08479532, -0.00657895,
        -0.00402047,  0.13596491, -0.01754386,  0.12426901,  0.08077485,
         0.15899123,  0.16959064, -0.10087719])
```

```
.base_values =
0.5657894736842105
```

```
.data =
array([ 1.2074225 ,  0.68849544, -0.95607147, -0.64519349, -1.31346187,
        -0.43414476,  0.93849623, -0.4464528 , -0.6815542 , -0.57895484,
         0.99874214, -0.70333816,  1.10304227])
```

"oldpeak" data and values

```
.values =
array([ 0.01023392, -0.06688596, -0.13450292,  0.08479532, -0.00657895,
        -0.00402047,  0.13596491, -0.01754386,  0.12426901,  0.08077485,
         0.15899123,  0.16959064, -0.10087719])
```

```
.base_values =
0.5657894736842105
```

```
.data =
array([ 1.2074225 ,  0.68849544, -0.95607147, -0.64519349, -1.31346187,
```

```

-0.43414476, 0.93849623, -0.4464528 , -0.6815542 , -0.57895484,
0.99874214, -0.70333816, 1.10304227])

"slp" data and values
.values =
array([ 0.01023392, -0.06688596, -0.13450292, 0.08479532, -0.00657895,
        -0.00402047, 0.13596491, -0.01754386, 0.12426901, 0.08077485,
        0.15899123, 0.16959064, -0.10087719])

.base_values =
0.5657894736842105

.data =
array([ 1.2074225 , 0.68849544, -0.95607147, -0.64519349, -1.31346187,
        -0.43414476, 0.93849623, -0.4464528 , -0.6815542 , -0.57895484,
        0.99874214, -0.70333816, 1.10304227])

"caa" data and values
.values =
array([ 0.01023392, -0.06688596, -0.13450292, 0.08479532, -0.00657895,
        -0.00402047, 0.13596491, -0.01754386, 0.12426901, 0.08077485,
        0.15899123, 0.16959064, -0.10087719])

.base_values =
0.5657894736842105

.data =
array([ 1.2074225 , 0.68849544, -0.95607147, -0.64519349, -1.31346187,
        -0.43414476, 0.93849623, -0.4464528 , -0.6815542 , -0.57895484,
        0.99874214, -0.70333816, 1.10304227])


"thall" data and values
.values =
array([ 0.01023392, -0.06688596, -0.13450292, 0.08479532, -0.00657895,
        -0.00402047, 0.13596491, -0.01754386, 0.12426901, 0.08077485,
        0.15899123, 0.16959064, -0.10087719])

.base_values =
0.5657894736842105


.data =
array([ 1.2074225 , 0.68849544, -0.95607147, -0.64519349, -1.31346187,
        -0.43414476, 0.93849623, -0.4464528 , -0.6815542 , -0.57895484,
        0.99874214, -0.70333816, 1.10304227])

```

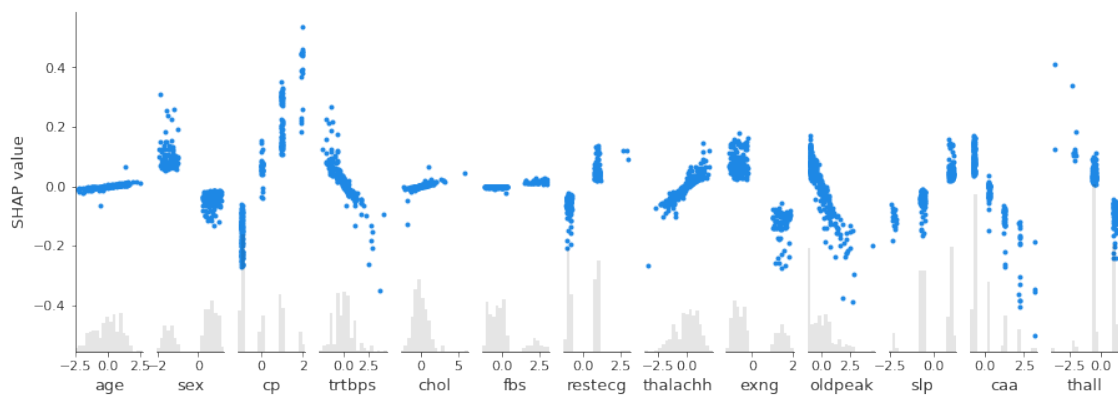
Printing Beeswarm SHAP Plot:





Printing Summary SHAP Bar Plot:







Printing Scatter SHAP Plot for All Features:







Printing Scatter SHAP Plot Between Features:





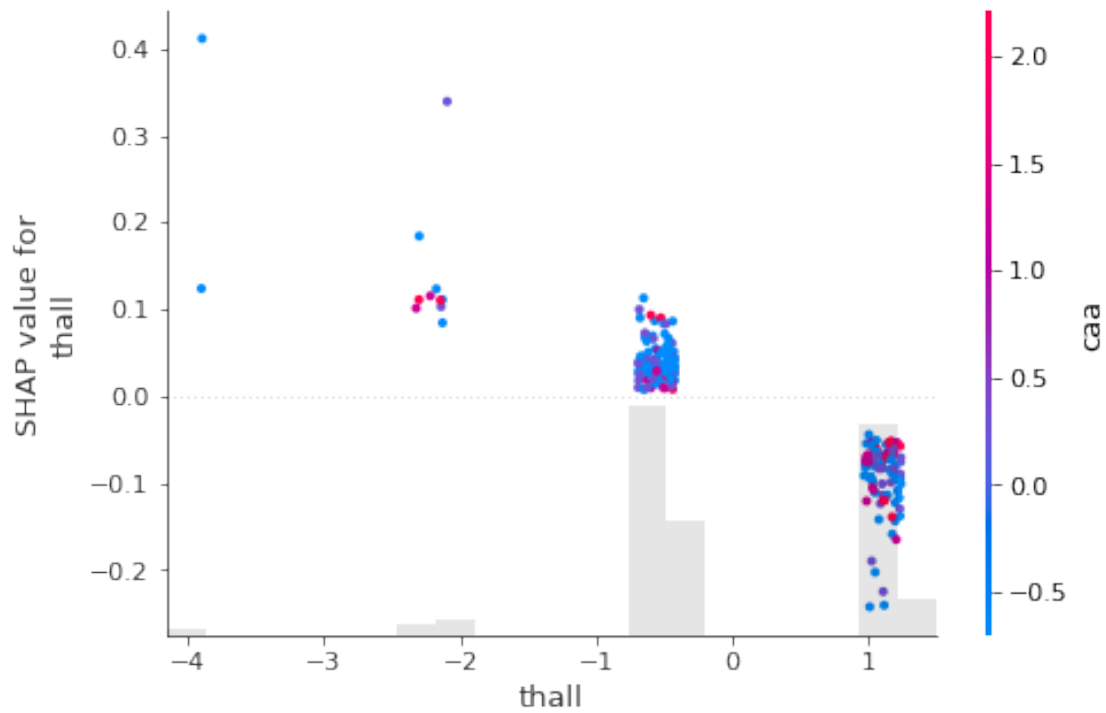




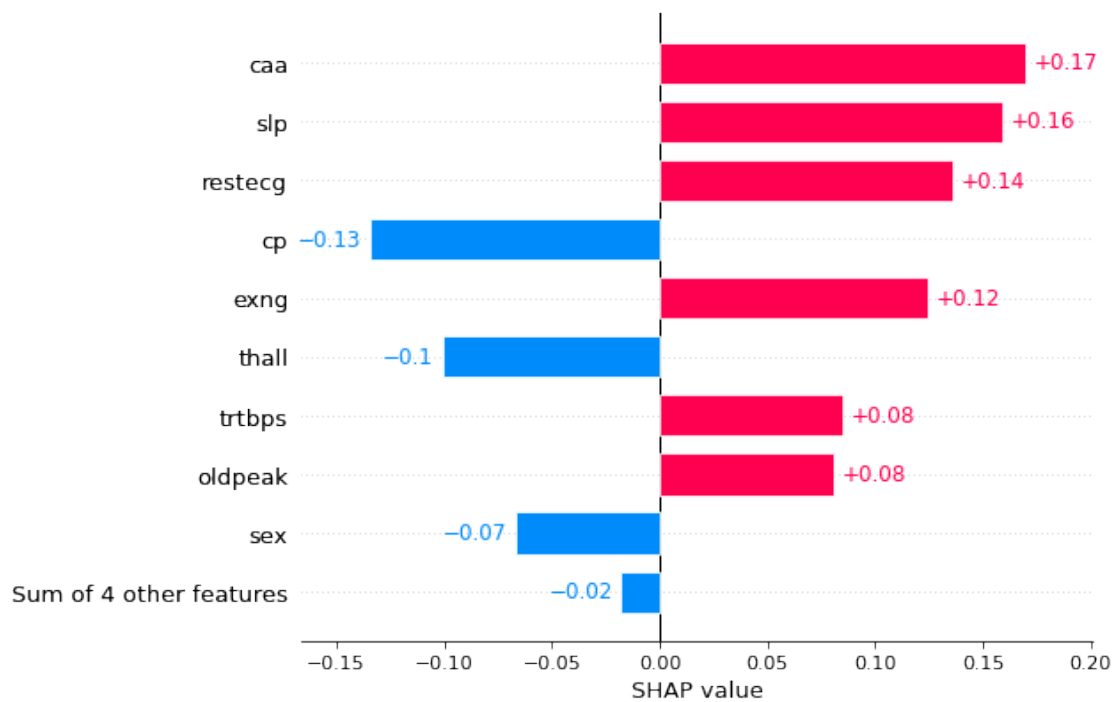




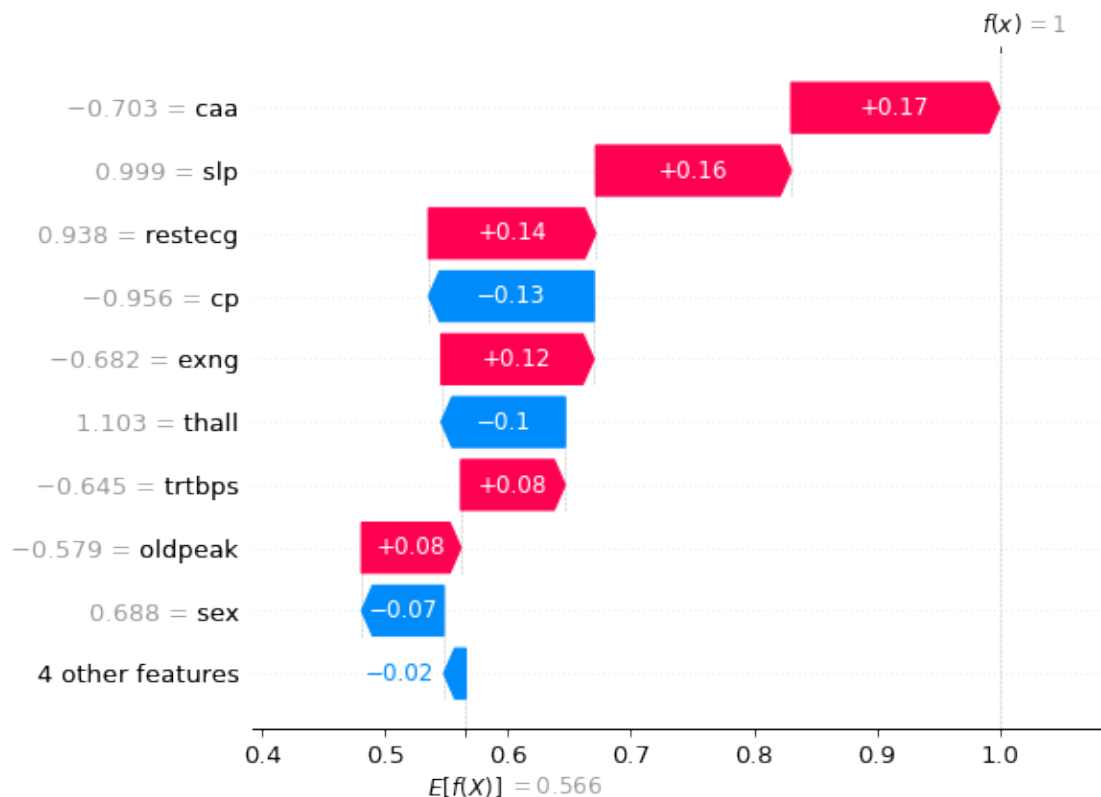




Printing Bar SHAP Plot for Local, Instance-Wise Effects:



Printing Waterfall SHAP Plot for Local, Instance-Wise Effects:



Printing Force SHAP Plot for Local, Instance-Wise Effects:

[11]: <shap.plots._force.AdditiveForceVisualizer at 0x12b874c10>

```
[12]: # References to put code together, compile, and test run of a Logistic
      ↪ Regression model using SHAP
      # Splitting CSV Into Train and Test Data by Nishank Sharma: https://
      ↪ medium.com/themlblog/splitting-csv-into-train-and-test-data-1407a063dd74
      # Logistic Regression in Python by Mirko Stojiljkovic: https://
      ↪ realpython.com/logistic-regression-python/
      # How to interpret machine learning (ML) models with SHAP values by
      ↪ Xiaoyou Wang: https://www.mage.ai/blog/
      ↪ how-to-interpret-explain-machine-learning-models-using-shap-values
```

```
#          SHAP explainer and models: https://www.mage.ai/blog/
→how-to-interpret-explain-machine-learning-models-using-shap-values
#          Dataset taken from: https://www.kaggle.com/datasets/
→rashikrahmanpritom/heart-attack-analysis-prediction-dataset?resource=download
#          scikit learn: https://scikit-learn.org/stable/modules/generated/
→sklearn.preprocessing.StandardScaler.html
#          SHAP Force Plots for Classification by Max Steele: https://medium.
→com/mlearning-ai/shap-force-plots-for-classification-d30be430e195
#          Machine Learning - Logistic Regression with Python by Nikhil
→Adithyan: https://medium.com/codex/
→machine-learning-logistic-regression-with-python-5ed4ded9d146
```